

# Maximum Likelihood Estimation Using RTMB

TMB through RTMB

[Click here to view presentation online](#)

Jim Bence

12 December 2024

## Outline: Part I

- Introductions
- Syllabus and our assumptions on background
- Course organization
- Philosophy and approach
- Expectations
- Break

## Course intro

- Welcome
- Housekeeping
  - Zoom, github, lecture recordings, communication, etc.
- Round-table introductions
  - Background, why are you here

## Syllabus

- Everything through Github
  - <https://github.com/QFCatMSU/MLE-Software/>
- Syllabus, presentations, code
- If you need help:
  - [bence@msu.edu](mailto:bence@msu.edu), [cahill11@msu.edu](mailto:cahill11@msu.edu)

## Course organization

- This is a hybrid course, with some participants in person and others online
- Each class session will be a mix of lecture and work on exercises
- For online participants, exercise work will be in breakout rooms
- No formal homework

## Our assumptions about your background

- Some previous experience with statistics including basic idea of fitting models to data
- Some previous use of R
- Some programming (understanding of functions, loops, conditional statements)
- Experience interpreting graphs
- You can get by without all this background, but you should expect to put in more time

## Jim's background

- BS in Biology University of Notre Dame
- PhD in Ecology and Masters in Statistics at Univ CA, Santa Barbara (both 1985).
- Worked on evaluation of environmental effects of San Onofre Nuclear Generating Station 1985-1989
- Mathematical Statistician (Stock Assessment Scientist) at NMFS-NOAA Tiburon Lab 1989-1994
- Faculty member at Michigan State University 1994- (retired from tenure stream position July 2023 - currently part time)

## Chris's background

- BS. University of Wisconsin Stevens Point (Fisheries science) 2011, MS. University of Alberta (Ecology/Evolutionary biology) 2014. PhD. University of Calgary (Ecology/Evolutionary biology) 2021
- Currently leading efforts to modernize and improve Great Lakes fisheries stock assessments.
- Expertise in hierarchical, state-space, and spatiotemporal modeling and quantitative tools used to inform resource management.
- Associate director of Quantitative Fisheries Center, MSU



## Our philosophy on statistical modeling

- Cookbook solutions rarely are adequate for real quantitative problems needed in ecology and resource management
- In class we will solve simple problems the hard way
  - This will make solving hard problems easier and position you to produce better solutions (Royle and Dorazio 2008)
- If you cannot write out your model you don't know what you did!
- Don't get lost in coding.
  - A good model you understand is critical.

## Software, implementation, website

- We are primarily going to use R and RTMB
- Recommend Rstudio
- Lectures and code will be available through GitHub
  - You do not need to know how to use GitHub, but that is where you can find code and presentations

---

---

---

---

## What is RTMB, why use it?

- R package for maximum likelihood fitting of arbitrarily complex models that incorporate random effects
- Nonlinear and non-normal models (within reason!)
- RTMB uses Template Model Builder (TMB) and TMB was inspired by AD Model Builder (ADMB)

## What is RTMB, why use it? A peak under the hood

- Automatic differentiation (as in ADMB)
- Laplace approximation to integrate out random effects (as ADMB)
- Automatic identification of parts of models that are connected (TMB/RTMB only)
- RTMB is much faster than ADMB for models with random effects
- No need for C++ coding (unlike ADMB)
- Major limitation – likelihood must be differentiable function of parameters and RTMB must “see” all the calculations in R

## Expectations

- Be kind
- Make an honest effort to learn this stuff
- Share your code
- This course lies at the intersection of mathematics, statistics, ecology, and numerical computing
  - Failure along the way is okay and to be expected and helps you learn and progress

## Some disclaimers

1. Maximum Likelihood Estimation is powerful but not without drawbacks and limitations
2. This is not a mathematical statistics course
3. We don't know everything
  - We will do my best to track down answers
4. Please ask questions

## Outline: Part II

- Overview of statistical modeling and MLE
- Building blocks
  - Probability, events, and outcomes
  - Random variables and random variates
    - Discrete versus continuous
  - Probability mass and density functions
  - Some common distributions
- The Likelihood function and maximum likelihood estimation

## A brief introduction to statistical modeling

- This class is about model-based inference
  - Focus on the development of arbitrarily abstract statistical models
- These models always contain:
  - Deterministic (i.e., systematic) components
  - Random (i.e., stochastic) components
- We estimate model parameters - can be of direct interest, or used to calculate something of interest
  - We are interested in uncertainty of estimates (SEs and CIs)



## Several example statistical models

- Regression model
- Hierarchical linear model
- Complex age structured assessment model

## Regression model example

$$y_i = f(\underline{\theta}, \underline{X}) + \varepsilon_i$$
$$L_i = L_{\infty} \left( 1 - e^{-K(a_i - t_0)} \right) + \varepsilon_i$$

---

---

Data: T. Brenden unpublished. Photo: E. Engbretson, USFWS <https://commons.wikimedia.org/w/index.php?curid=3720748>

## Hierarchical linear model example

- Weight is power function of length multiplied by error
  - On log scale the relationship is linear with additive error
- $i$  represents ponds,  $j$  fish within ponds
- Intercept and slope vary randomly among ponds, residual variance is pond specific
- $Y \sim N(a, b)$  means  $Y$  is normal, with mean  $a$  and variance  $b$

$$\log W_{ij} = a_i + b_i \log(L_{ij}) + \varepsilon_{ij}$$

$$a_i \sim N(\alpha, \sigma_a^2), b_i \sim N(\beta, \sigma_b^2), \varepsilon_{ij} \sim N(0, \sigma_i^2)$$

## State Space Catch at age model

$$\log N_{4,y} = \log N_{4,y-1} +$$

$$\varepsilon_y^{(R)}; \varepsilon_y^{(R)} \sim N(0, \sigma_R^2)$$

$$\log N_{a,1986} = \log N_{4,1986-(a-4)} -$$

$$\sum_4^{a-1} \log \bar{Z}_a, 4 < a \leq 9$$

$$\log N_{a,1986} = 0, a > 9$$

$$\log N_{a,y} = \log N_{a-1,y-1} - Z_{a-1,y-1},$$

$$4 \leq a < A, \log N_{A,y} =$$

$$\log(N_{A-1,y-1} e^{-Z_{A-1,y-1}} + N_{A,y-1} e^{-Z_{A,y-1}})$$

$$Z_{a,y} = M + \sum_{G=g,t} F_{a,y,G}$$

## State Space Catch at age model (continued)

$$F_{a,y,G} = q_{a,y,G} E_{y,G}, G = g, t$$

$$\log q_{y,G} = \log q_{y-1,G} + \varepsilon_y^{(G)}; \varepsilon_y^{(G)} \sim$$

$$N(0, \Sigma_G), G = g, t$$

$$\Sigma_{a,\tilde{a}} = \rho^{|a-\tilde{a}|} \sigma_a \sigma_{\tilde{a}}, 4 < a \leq A,$$

$$4 < \tilde{a} \leq A$$

$$B_y^{(Spawn)} = \sum_{a=4}^A m_{a,y} W_{a,y}^{(spawn)} \log N_{a,y}$$

$$C_{a,y,G} = \frac{F_{a,y,G}}{Z_{a,y}} N_{a,y} (1 - \exp(-Z_{a,y})),$$

## Probability

- Whole books about definitions and meaning.
- I follow a frequentist definition for intuition, while recognizing that there is some logic to Bayesian claims of degree of belief
  - Frequentist definition: The long run proportion of of times an event occurs under identical conditions
  - Statisticians sometimes distinguish outcomes from events. Outcomes are really elementary events. Event might be catching a fish in 7 to 8 inch bin, outcome would be catch a fish and measure its length.

## Basic properties (axioms) of probabilities

- The sum of probabilities over all possible mutually exclusive events is 1.0 (So something will happen)
- Probability of any given event is  $\geq 0$  (and  $\leq 1$ )
- The probability of the union of mutually exclusive events is the sum of their separate probabilities
  - $P(A \cup B) = P(A) + P(B) - P(A \cap B)$
- if A and B independent  $P(A \cap B) = P(A)P(B)$

## Conditional probability

$$P(A \mid B) = \frac{P(A \cap B)}{P(B)}$$

- “|” read as “given or conditional on – Probability of A given B
- Conditional probabilities recognize that the occurrence of event B can provide information on whether event A will occur
- Convince yourself that  $P(A \mid B) = P(A)$  if A and B independent



## Conditional probability

---

[stats.stackexchange.com/questions/587109](https://stats.stackexchange.com/questions/587109)

## Random variables (words)

- Technical definition is that they are functions that convert probability spaces for events/outcomes to numeric results.
  - Ironically they are neither random nor variables!
- Less technically (but still techno speak!) they describe the numeric outcome of a random process. I.e., they are not a number (or vector/matrix of numbers) but rather the process of producing them.
- A random variate is a particular numeric outcome
- Text books say usually capital letters used for random variables and lower case for random variates.

## Random variables: math expression for simple example (coin flip)

WE FLIP A COIN AND CALL A HEADS **1** AND A TAILS **0**:

$$\text{suppose } p = \Pr(Y = 1)$$

$$\Pr(Y = 1) + \Pr(Y = 0) = 1$$

$$\Pr(Y = 0) = 1 - p$$

- $\Pr$  = probability
- We could say the random variable  $Y$  has a Bernoulli distribution

## Bernoulli probability mass function (pmf)

$$\Pr(Y = y) = p^y(1 - p)^{1-y}$$

- The pmf is a function that calculates the probability given the random variate ( $y$  value) and the parameter(s) (here  $p$ )  
 $y$  is observed datum
  - pmf for discrete outcomes
- If this was a continuously distributed random variable we would use probability *density* function (pdf)

## pmf and pdf notation

- A conventional notation for this stuff is  $f(y \mid \theta)$
- Sometimes with subscript for random variable:  $f_Y(y \mid \theta)$
- Conditional bit indicates that the probability of an observed value  $y$  depends on parameter(s)  $\theta$  used to specify the distribution of the random variable  $Y$
- Notation for Benouilli random variable:  $f(y \mid \theta) = p^y(1 - p)^{1-y}$

## Some common statistical distributions




Figure Reference

## More notation notes for everyone's sanity

- In general we will provide the pmf (or pdf) expressed as a function of  $y$  and the parameters of the distribution.
- For example, will use  $y \sim N(\mu, \sigma^2)$  to indicate a random variable  $Y$  is normally distributed with mean  $\mu$  and variance  $\sigma^2$
- In general regular font for scalars, bold for vectors and matrices

## More notation

- You might see it this way too:

$$\text{Normal}(y \mid \mu, \sigma^2) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2} \left(\frac{y - \mu}{\sigma}\right)^2\right).$$



## Discrete versus Continuous random variables

- Discrete means the set of possible outcomes is countable with each possible value having an associated probability (calculable from the pmf).
- Continuous means not countable (generally this means there are infinite numbers of possible values between any two other possible values).  $\Pr(y)$  for any particular  $y$  is 0. So we use a probability density function.
- Intuition/common sense sometimes used to choose between the two. E.g., catch or CPUE often modeled as continuous

## Probability density function

---

- $\Pr(c_1)=\Pr(c_2)=0$  !
- Area under the pdf function gives probability for interval
- $\Pr(c_1 < x < c_2)=\Pr(x < c_2)-\Pr(x < c_1)$

## Cumulative distribution function

- $F(x) = \Pr(X \leq x)$
- For continuous variables, the derivative of  $F(x)$  with respect to  $x$  is  $f(x)$  (the density)
  - Why? Does this make sense?

## Joint probability density and mass functions

Vector of observed values, with elements having the same pdf/pmf or different ones and these random variables might be independent or not:

$$f(x_1, x_2, \dots, x_k) = f(\mathbf{x})$$

Special case of each element representing an independent random variable:

$$f(\mathbf{x}) = f_{X_1}(x_1)f_{X_2}(x_2)\dots f_{X_k}(x_k)$$

Special special case of independent and identically distributed (iid) random variables:  $f(\mathbf{x}) = f(x_1 | \theta_1)f(x_2 | \theta_2)\dots f(x_k | \theta_k) = \prod_{i=1}^k f(x_i | \theta_i)$

These special cases very important for practical MLE work!

## The likelihood function

- No new math!!!
- The likelihood function is just the joint pdf re-expressed as a function of the parameters:  $f(\theta|\mathbf{x})$

## Maximum likelihood estimation

- Adjust  $\theta$  until  $f(\theta|\mathbf{x})$  is maximized
- The rest is “just” details :->

## A numerical example of joint likelihood

- $\mathbf{x} = \{10.72, 7.23, 10.07, 8.62, 8.55\}$
- Each observation ( $x$ ) independent from a common normal distribution with mean 10, variance 2 (i.e., they are iid)
- Calculate the likelihood of these data, i.e.,  $f(10.72)f(7.23)f(10.07)f(8.62)f(8.55)$  using R
- Hints. The result is just a single number. You can calculate the pdf of  $f(x)$  for a normal distribution in R using the `dnorm` function. The `dnorm` function uses sigma (SD) not sigma squared (variance)
- Time permitting - generalize your solution as a R function to calculate likelihood for any vector  $x$ , with specified mean, variance.

## Working with the log likelihood (preferred for numerical reasons)

- Perhaps obviously, if you adjust parameters to maximize the log of the likelihood function this will also maximize the likelihood.
- RTMB and most software minimizes the negative log likelihood rather than maximizing log likelihood (convention)
- Working on the log-scale improves numerical performance.
  - The joint log likelihood for independent observations is the SUM (rather than product) of the individual  $\log f(x)$  values



## Negative log likelihood versus $\mu$ for five iid observations from a Normal distribution (known variance of 2)

---

NLL versus  $\mu$

## Detection probability example - Sneak turtles

- 30 turtles released in pond two recaptured later
- Release is trial, recapture is a success
- If each trial is independent, number of successes binomial
- Just one observation -  $y$  successes. Likelihood just the pdf

$$p(y \mid \theta) = \frac{n!}{y!(n - y)!} \theta^y (1 - \theta)^{n-y}$$

## Finding the MLE - The sneaky example

- Determine the detection probability of invasive Alabama sneak turtles by releasing some into a small pond and counting them on a later occasion
- We release  $n = 30$  and later count  $y = 2$  sneakers
  - successes vs. trials
- What's probability of detecting a turtle, i.e.,  $\theta$ ?

$$p(y \mid \theta) = \frac{n!}{y!(n - y)!} \theta^y (1 - \theta)^{n-y}$$

## Methods for finding MLEs

- Analytical solution (involves derivatives)
- Grid search
- Iterative searches
  - Non-derivative methods
  - Derivative methods (such as quasi-Newton)

## Role of derivatives in finding MLEs

NLL as function of single parameter with derivatives

- Derivatives of NLL with respect to parameters zero at minimum
- Second derivatives of NLL with respect to parameters are positive at minimum

## Sneak turtle analytical approach

$$NLL = -y \log \theta - (n - y) \log(1 - \theta) + C$$

$$\frac{\partial NLL}{\partial \theta} = \frac{n-y}{1-\theta} - \frac{y}{\theta} = 0$$

$$\theta = \frac{y}{n}$$

## Grid search for sneak turtle example

```
1 y <- 2 # successes
2 n <- 30 # trials
3 eps <- 1e-6 # buffer because theta must lie within [0,1]
4 theta <- seq(eps, 1 - eps, length.out = 1e4) # sequence to thetas to
5 # plug thetas into pmf
6 logLike=y * log(theta) + (n - y) * log(1 - theta)
7 MLE <- theta[which.max(logLike)] # maximum likelihood estimate
8 my_data <- data.frame(logLike, theta) # for plotting if you want
```

## Analytical vs. grid search MLEs

```
1 y / n # analytical theta MLE
```

```
[1] 0.06666667
```

```
1 MLE # grid search approximation
```

```
[1] 0.06670754
```



## Plot of log likelihood function for sneak turtle

```
1 curve(y * log(x) + (n - y) * log(1 - x), from = eps, to = 0.3, xlab = "theta")
```

---

## Exercise - use built in binomial distribution

- Calculate the log likelihood at the MLE using `logLike = dbinom(y,n,MLE,log=TRUE)`
- Compare with `logLike = y * log(MLE) + (n - y) * log(1 - MLE)`
  - Why are they different?
- Steal the grid search code, run it.
  - Then replace `logLike = y * log(theta) + (n - y) * log(1 - theta)` by `logLike=dbinom(y,n,theta,log=TRUE)`
  - You should get the same grid search estimate

## Another exercise or interest of time demo

- For the previous sample of five observations from normal find the MLE estimate of the mean assuming the variance known equal to 2 by conducting a grid search
- Think about how you would do grid search to find MLE estimates of both the mean and variance at the same time

## The regression case

- Observations assumed independent but not identically distributed.
- The mean varies among observations and in this simple case variance is same for all observations:  $y_i \sim N(\mu_i, \sigma^2)$
- Cannot estimate mean as a parameter for every observation
- But can calculate it as function of parameters, e.g.,  $\mu_i = \alpha + \beta X_i$
- general message estimated parameters are not the same as the distributional parameters for the pdfs/pmfs
- What are the estimated (model) parameters?

## Pseudocode for the regression problem

- Specify  $\alpha$ ,  $\beta$ , and  $\sigma^2$
- Calculate  $\mu_i$
- Calculate NLL

Search over different values of  $\alpha$ ,  $\beta$ , and  $\sigma^2$  and repeat 1-3 until you find the values that minimize the NLL

## Two ways to frame the regression model

$$y_i = \mu_i + \epsilon_i = \alpha + \beta * X_i + \epsilon_i$$

$$\epsilon \stackrel{iid}{\sim} N(0, \sigma^2)$$

- Previously we modeled mean, which was a distributional parameter. Now we write a model for the individual observations. We can write the likelihood in terms of the errors, or the observations - The two are theoretically equivalent!
- Note but don't worrying about. The standard model notation uses random quantities on the RHS but specifies the random variate on the LHS

## Properties of MLEs

- Terminology: ML Estimator versus ML Estimate
- Ideal estimator is lowest variance among unbiased estimators
- MLE not guaranteed to do this!
- MLEs are consistent, meaning estimates will become closer to correct values and sample sizes increase
  - Asymptotically unbiased
  - Errors get smaller with more data

## Familiar example of bias for an MLE

- MLE for variance of normal random sample:  $\hat{\sigma}^2 = \sum (x_i - \hat{\mu})^2 / k$
- Expected value:  $E(\hat{\sigma}^2) = \frac{k-1}{k} \sigma^2$
- Standard (unbiased) estimator:  $\hat{\sigma}_u^2 = \sum (x_i - \hat{\mu})^2 / (k - 1)$